

Operating Systems

Memory Stress Tests in CGROUPS

Ameet Deshpande - (CS15B001)

October 31, 2017

1 What are CGROUPS

- CGROUPS are short for Control Groups which was introduced in Linux. Cgroups basically allow us to allocate resources such as CPU time, system memory, network bandwidth among user-defined groups of tasks. This seems very advantageous when we are expecting some kind of workload for the system.
- One potential downside is the fact that Control Groups are still user defined. If there was some way to automatically cluster process to enhance scheduling (like GAS) we could place those processes in the Control Groups. But nevertheless, Linux gives us the power to group processes and enforce certain constraints on the whole set of processes rather than single processes.
- CGROUPS also have a hierarchy like processes in Linux. Child CGROUPS can inherit attributes from the parent CGROUPS. One difference however is that in Linux Processes, all of them are children of *init*. In other words, it is a single Hierarchy. In CGROUPS however, there are multiple hierarchies.

2 Applying Stress Tests

- The assignment mainly requires us to enforce certain memory conditions on the CGROUPS. We use the package called *stress* which is available for enforcing the memory constraints.
- *stress* spawns processes which stress the hardware and OS according to the parameters provided.
- The command *stress -m 3* spawns 3 processes which stress the hardware with 256MB each. We can further tune the amount of memory each process occupies.
- The command *stress -m 4 -vm-bytes 300M* basically stresses the hardware with 4 processes of 300MB each. These commands thus give us a lot of control over how much memory each process is using. At the same time, it allows us to get multiple processes under the same CGROUP.

3 Applying memory constraints to CGROUPS

- CGROUPS can be created and constraints can be applied to them. The following commands illustrate the same.

```
1 sudo cgcreate -g memory:osgroup
2 sudo cgset -r memory.limit_in_bytes=$((25*1024*1024))$ osgroup
3 sudo cgexec -g memory:osgroup stress -m 3
```

- The first command creates a control group. The second command sets the memory limit and the third spawns processes under the control group. *htop* is being used for monitoring the processes.

4 Experiments

4.1 Part 1

- *stress* basically provides processes which run for a very long time. We thus use them to test how the memory is doing. We set the CGROUPS memory limit as 100M.
- We spawn 2 processes with 25M each, with a memory limit of 100M. We expect the programs to run without a problem and the screenshot below suggest the same.

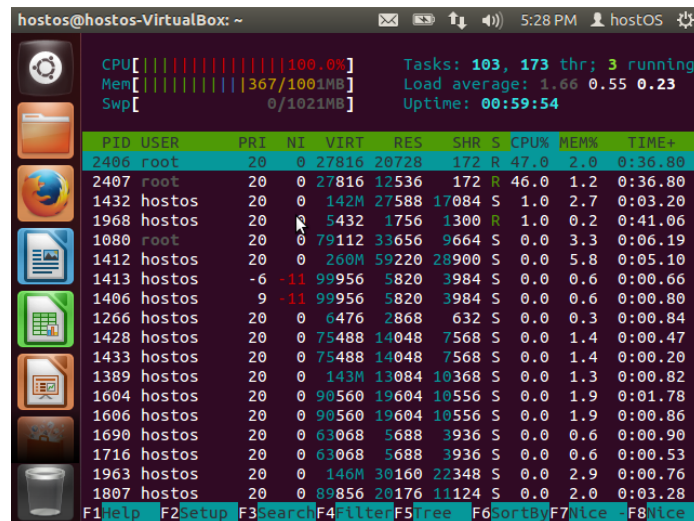


Figure 1

- Increasing the number of processes to 4, we can see that the memory utilization increased.




Figure 2

- Increasing the number of processes forces the swap space to be used as can be seen below. It can be seen that about 140M of SWAP is being used along with the main memory. We expect such programs to run slower given that large amount of disk I/O will be involved.

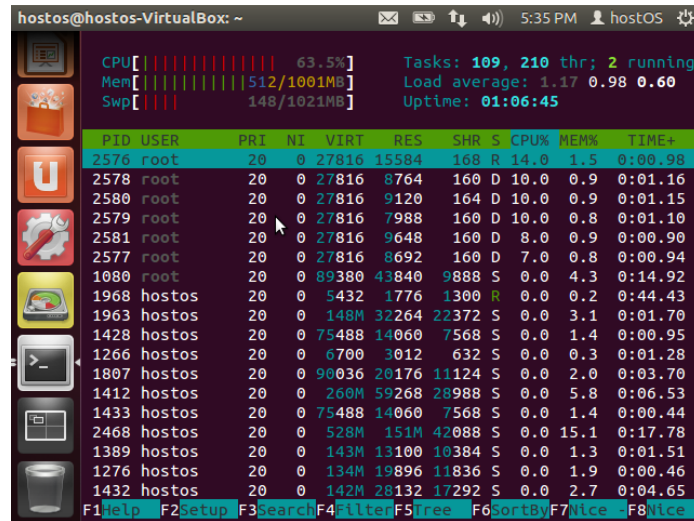


Figure 3

4.2 Part 2

- This section aims to check how the swappiness of the system effects the processes. We set the swappiness of the system to 0 to check how the processes respond. We expect the process to not even run when the amount of memory is not enough for it. The following is the result.




Figure 4: No Swap

4.3 Performance

- To check how the performance of the program varies with the memory limit, a small Java program was written. The code is attached below.

```

1 import java.util.*;
2
3 public class test {
4     public static void main(String[] args){
5         LinkedList<String> a = new LinkedList<String>();
6         String temp = "oslab_projects";
7
8         for(int i=0;i<1024*1024*8;++i){
9             a.add(temp);
10        }
11
12        System.out.println("Done");
13    }
14 }

```

The CGROUP size was limited to just 50MB, while this program takes about 100MB. The JAVA heap size was made around 1000MB so that garbage collector does not interfere in the performance measurement. The following screenshot measure the performance of the java programs with and without the CGROUP.

```

hostos@hostos-VirtualBox: ~/Downloads
hostos@hostos-VirtualBox:~/Downloads$ time java test
Done
real    0m0.811s
user    0m0.756s
sys     0m0.028s
hostos@hostos-VirtualBox:~/Downloads$ time sudo cgexec -g memory:osgroup java test
[sudo] password for hostos:
Done
real    0m8.900s
user    0m0.944s
sys     0m1.228s
hostos@hostos-VirtualBox:~/Downloads$

```

Figure 5: No Swap

- As can be seen in the screenshot, when the heap size is kept unbounded, the program takes only 0.8 seconds to execute. When the Memory of the CGROUP is limited, the program takes 8.9 seconds to run. We can clearly see how usage of the SWAP affects the program time.

5 Soft and Hard Limits

- There are two ways of handling limits in CGROUPS. If the hard limit is selected, the total amount of physical memory used by the sum of all processes in the job will not be allowed to exceed the limit. If the processes try to allocate more memory, the allocation will succeed, and virtual memory will be allocated, but no additional physical memory will be allocated. The system will keep the amount of physical memory constant by swapping some page from that job out of memory. If the soft limit is selected, the job will be allowed to go over the limit if there is free memory available on the system. Only when there is contention between other processes for physical memory will the system force physical memory into swap and push the physical memory used towards the assigned limit. The following command helps us create CGROUPS with soft limits.

```
1 memory.soft_limit_in_bytes
```

- We perform a very interesting experiment. We create 2 CGROUPS with the same limits and same tasks. We assign one of them Soft Limits and the other Hard Limits. We keep the memory usage of the processes such that a crash is expected. We see which of the processes crash first.




Figure 6: Fail




Figure 7: Success

- The results are as expected. Given that the first group has hard-limits, the processes get killed. But the second CGROUP has soft limits and hence it continues executing.