

Printing all subdirectories and files in current directory

Ameet Deshpande (CS15B001)

October 23, 2017

1 Exploring Stat Structure

- *stat* is a system call in C that can be used to get information about files using their file names. It gets its information from a structure which also contains the Inode number. Once the Inode Number is obtained, all information about the file can be accessed. The following is the structure of *stat*.

```
1 struct stat {
2     dev_t      st_dev;      /* ID of device containing file */
3     ino_t      st_ino;      /* inode number */
4     mode_t     st_mode;     /* protection */
5     nlink_t    st_nlink;    /* number of hard links */
6     uid_t      st_uid;      /* user ID of owner */
7     gid_t      st_gid;      /* group ID of owner */
8     dev_t      st_rdev;     /* device ID (if special file) */
9     off_t      st_size;     /* total size, in bytes */
10    blksize_t   st_blksize;  /* blocksize for file system I/O */
11    blkcnt_t    st_blocks;   /* number of 512B blocks allocated */
12    time_t      st_atime;    /* time of last access */
13    time_t      st_mtime;    /* time of last modification */
14    time_t      st_ctime;    /* time of last status change */
15    };
```

As can be seen, the *st_ino* field gives the inode number of the file. The following code is written to illustrate the use of the *stat* structure.

```
1 #include <sys/stat.h>
2 #include <stdio.h>
3 #include <time.h>
4
5 int main(){}
6     struct stat s;          // Structure to store file information
7     stat("os.txt", &s);    // Store the information about file
```

```

8      // Print details about the file
9      printf("INode Number of file: %ld\n", (long) s.st_ino);
10     printf("Size of the file is: %ld KB\n", (long) s.st_size);
11     printf("Last accessed %s", ctime(&s.st_atime));
12 }

```

Note the usage of *st_ino* to print the inode number and *st_size* to print the size of the file names os.txt. (Code is attached)

The following output is produced when the file is run.

```

1 INode Number of file: 58448
2 Size of the file is: 3 KB

```

2 Lising the Directories

- The standard library files in *dirent.h* allow us to read directories. It provides functions *readdir*, *opendir* and *closedir* which can be used as utility functions to read, open and close files. One important catch when we are recursively printing the files and directories inside the current directory is to not include *.* and *...*. This is because it will cause a stack overflow as the current directory calls the current directory and the parent directory. There is a statement which checks if the directory name are any of those to avoid this. The program that is turned in basically does a walk over all the directories and files in the current directory and prints the inode number along with it. The walk is done using a *DIR* pointer which stores the opened file's information. After getting the *DIR* pointer, *DT_DIR* flag is used to check if it is a directory. If it is a directory, a recursive call is made. If it is a file, the information about the file is printed along with the inode. Note that this is a very general procedure, as once the *stat* structure is obtained for a file, the inode number can be obtained. The inode number contains all the information/metadata with respect to the file. The code is attached below.

```

1 #include <unistd.h>
2 #include <sys/types.h>
3 #include <dirent.h>
4 #include <stdio.h>
5 #include <string.h>
6 #include <sys/stat.h>
7
8 void list_recursive (const char *name, int spaces){
9     DIR *dir;
10    struct dirent *entry;
11
12    if (!(dir = opendir(name))) // If the directory does not exist {
13        return;

```

```

14     }
15
16     while((entry = readdir(dir)) != NULL){
17         if(entry->d_type == DT_DIR) // If the "file" is a directory, call
            recursively
18         {
19             char path[1024];
20             if(strcmp(entry->d_name, ".") == 0 || strcmp(entry->d_name,
                "..") == 0) // If parent directory or current, then dont
                print
21             {
22                 // Else
23                 stack overflow
24                 continue;
25             }
26             snprintf(path, sizeof(path), "%s/%s", name, entry->d_name);
27             // Print the path to a string
28             struct stat s;
29             stat(path, &s);
30             printf("%ld :: ", (long) s.st_ino);
31             printf("%s[%s]\n", spaces, "", entry->d_name);
32             list_recursive (path, spaces + 2);
33         }
34     else{
35         struct stat s;
36         char path[1024];
37         snprintf(path, sizeof(path), "%s/%s", name, entry->d_name);
38         stat(path, &s);
39         printf("%ld :: ", (long) s.st_ino);
40         printf("%s- %s\n", spaces, "", entry->d_name);
41     }
42 }
43
44 int main(void) {
45     list_recursive (".", 0); // Call the function on current directory
46     return 0;
47 }

```

The *snprintf* function is a very helpful utility function which prints to a string what it would have otherwise printed to *stdout*.

- Sample output will look like the following.

```

1 58393 :: - a.out
2 58448 :: - os.txt

```

```
3 58408 :: - print_inode_number.c
4 58501 :: - print_recursive.c
5 58502 :: [random1]
6 58507 :: - file.c
7 58390 :: - file2.txt
8 58503 :: [random2]
9 58506 :: - file.c
10 58504 :: [random3]
11 58505 :: - file3.c
```

3 Exploring the inode structure

The Inode structure source code is available [here](#). It is clear from the same that the stat structure has to use this behind the hood. A few lines from the structure are pasted below.

```
1  loff_t          i_size ;
2  struct timespec    i_atime;
3  struct timespec    i_mtime;
4  struct timespec    i_ctime;
```

The stat structure gets its information from the above fields that are present in the inode structure.

```
1  union {
2      struct hlist_head i_dentry;
3      struct rcu_head    i_rcu;
4  };
```

From the above, it can be seen that the first field contains the directory entry of the inode. The structure also has fields for locks and semaphores which can be used while reading from or writing to the blocks. The DIR pointers and the Stat structures can thus get all the information needed from the inode structure. The whole structure is attached below for reference. Hardlinks are used in the directory structure. It is used to associate a name with a file on the file system. Note that the name means nothing to the operating system. It is just a way to project the information to the user. Under the hood, everything is implemented using inodes as they contain all the information that is required to read all the files and sub directories.

4 Large Files and Inodes

Linux file system is based on the fact that the OS is made of a large number of files of small size. The inode thus points to direct data blocks which have enough space to store files of small size, like *4 KB*. When the file sizes become larger, the indirect blocks have to be used. The indirect blocks basically have pointers which point

to other blocks on disk. There are several levels of indirect blocks. The second level indirect blocks will thus have pointers to blocks which have pointers to the actual data blocks. Thus, it can clearly be seen that in a normal unix file system, accessing large files will be cumbersome and will take a large amount of time compared to smaller files. For small files, the inode entries for indirect blocks will be empty. However as the file becomes larger, there wont be enough space to store the file in direct blocks. The indirect blocks will thus store an address of a block which stores the final data. The advantage is that the file does not have to be stored in contiguous locations and there will be lesser problems in file allocations.

5 EXT4 File System

- This file system was designed to handle larger files which will suffer in the file system described in the preceding section. How it does the same is to use the concept of extents. This basically reduces the fragmentation that the other file systems introduce. Extents can be used to map physical disk sizes of about $128MiB$, using just $4KiB$ of continuous space. The inode in ext4 can store upto 4 extents. When the number of extents required becomes more than 4, extents are indexed by a tree. This makes it more efficient to access data. The file system *XFS* is recommended when even larger files need to be stored. In EXT3 every-time an append operation is called, a new block allocator was called. When concurrent processes are writing to the same files, all of them are given different blocks and this causes the issue of fragmentation. This is avoided in EXT4 by delaying the allocation of blocks and thus using the same block to write. EXT4 is thus more suited for Concurrent Programming.

6 Inode Struct

```

1 struct inode {
2     umode_t      i_mode;
3     unsigned short i_opflags ;
4     kuid_t       i_uid ;
5     kgid_t       i_gid ;
6     unsigned int  i_flags ;
7
8     #ifdef CONFIG_FS_POSIX_ACL
9         struct posix_acl *i_acl ;
10        struct posix_acl *i_default_acl ;
11    #endif
12
13    const struct inode_operations *i_op;
14    struct super_block *i_sb;
15    struct address_space *i_mapping;
```

```

16
17 #ifndef CONFIG_SECURITY
18     void          *i_security ;
19 #endif
20
21     /* Stat data, not accessed from path walking */
22     unsigned long   i_lino ;
23     /*
24      * Filesystems may only read i_nlink directly .  They shall use the
25      * following functions for modification:
26      *
27      * (set | clear | inc | drop)_nlink
28      * inode_(inc | dec)_link_count
29      */
30     union {
31         const unsigned int i_nlink;
32         unsigned int __i_nlink;
33     };
34     dev_t          i_rdev ;
35     loff_t          i_size ;
36     struct timespec i_atime;
37     struct timespec i_mtime;
38     struct timespec i_ctime;
39     spinlock_t      i_lock ; /* i_blocks , i_bytes , maybe i_size */
40     unsigned short   i_bytes ;
41     unsigned int      i_blkbits ;
42     enum rw_hint      i_write_hint ;
43     blkcnt_t         i_blocks ;
44
45 #ifdef __NEED_I_SIZE_ORDERED
46     seqcount_t        i_size_seqcount ;
47 #endif
48
49     /* Misc */
50     unsigned long      i_state ;
51     struct rw_semaphore i_rwsem;
52
53     unsigned long      dirtied_when; /* jiffies of first dirtying */
54     unsigned long      dirtied_time_when;
55
56     struct hlist_node i_hash;
57     struct list_head   i_io_list ; /* backing dev IO list */
58 #ifdef CONFIG_CGROUP_WRITEBACK
59     struct bdi_writeback *i_wb; /* the associated cgroup wb */
60

```

```

61  /* foreign inode detection, see wbc_detach_inode() */
62  int      i_wb_frn_winner;
63  u16      i_wb_frn_avg_time;
64  u16      i_wb_frn_history;
65  #endif
66  struct list_head i_lru;      /* inode LRU list */
67  struct list_head i_sb_list;
68  struct list_head i_wb_list; /* backing dev writeback list */
69  union {
70      struct hlist_head i_dentry;
71      struct rcu_head i_rcu;
72  };
73  u64      i_version;
74  atomic_t i_count;
75  atomic_t i_dio_count;
76  atomic_t i_writecount;
77  #ifdef CONFIG_IMA
78  atomic_t i_readcount; /* struct files open RO */
79  #endif
80  const struct file_operations *i_fop; /* former ->i_op->default_file_ops */
81  struct file_lock_context *i_flctx;
82  struct address_space i_data;
83  struct list_head i_devices;
84  union {
85      struct pipe_inode_info *i_pipe;
86      struct block_device *i_bdev;
87      struct cdev *i_cdev;
88      char *i_link;
89      unsigned i_dir_seq;
90  };
91
92  __u32      i_generation;
93
94  #ifdef CONFIG_FSNOTIFY
95  __u32      i_fsnotify_mask; /* all events this inode cares about */
96  struct fsnotify_mark_connector __rcu *i_fsnotify_marks;
97  #endif
98
99  #if IS_ENABLED(CONFIG_FS_ENCRYPTION)
100  struct fscrypt_info *i_crypt_info;
101  #endif
102
103  void      *i_private; /* fs or device private pointer */
104  } __randomize_layout;

```

7 References

- Wikipedia
- Stack Overflow
- [GitHub Repo](#)
- [Stat Structure](#)
- Free-Electrons